

```
In [ ]: Loan Approval_ML milestone_project
```

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    roc_auc_score, roc_curve, confusion_matrix
)
```

```
In [29]: df= pd.read_csv("loan_approval_dataset.csv")
df
```

	loan_id	no_of_dependents	education	self_employed	income annum	loan_amount	loan_term	cibil_score	residential_assets_value	commercial_assets_value	luxury_assets_value	bank_asset_value	loan_status
0	1	2	Graduate	No	9600000	29900000	12	778	2400000	17600000	22700000	8000000	Approved
1	2	0	Not Graduate	Yes	4100000	12200000	8	417	2700000	2200000	8800000	3300000	Rejected
2	3	3	Graduate	No	9100000	29700000	20	506	7100000	4500000	33300000	12800000	Rejected
3	4	3	Graduate	No	8200000	30700000	8	467	18200000	3300000	23300000	7900000	Rejected
4	5	5	Not Graduate	Yes	9800000	24200000	20	382	12400000	8200000	29400000	5000000	Rejected
...
4264	4265	5	Graduate	Yes	1000000	2300000	12	317	2800000	500000	3300000	800000	Rejected
4265	4266	0	Not Graduate	Yes	3300000	11300000	20	559	4200000	2900000	11000000	1900000	Approved
4266	4267	2	Not Graduate	No	6500000	23900000	18	457	1200000	12400000	18100000	7300000	Rejected
4267	4268	1	Not Graduate	No	4100000	12800000	8	780	8200000	700000	14100000	5800000	Approved
4268	4269	1	Graduate	No	9200000	29700000	10	607	17800000	11800000	35700000	12000000	Approved

4269 rows × 13 columns

```
In [30]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4269 entries, 0 to 4268
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   loan_id          4269 non-null    int64  
 1   no_of_dependents 4269 non-null    int64  
 2   education         4269 non-null    object  
 3   self_employed     4269 non-null    object  
 4   income_annum     4269 non-null    int64  
 5   loan_amount       4269 non-null    int64  
 6   loan_term         4269 non-null    int64  
 7   cibil_score       4269 non-null    int64  
 8   residential_assets_value 4269 non-null    int64  
 9   commercial_assets_value 4269 non-null    int64  
 10  luxury_assets_value 4269 non-null    int64  
 11  bank_asset_value  4269 non-null    int64  
 12  loan_status       4269 non-null    object  
dtypes: int64(10), object(3)
memory usage: 433.7+ KB
```

```
In [31]: for col in df.columns:
    print(repr(col))
```

```
'loan_id'
'no_of_dependents'
'education'
'self_employed'
'income_annum'
'loan_amount'
'loan_term'
'cibil_score'
'residential_assets_value'
'commercial_assets_value'
'luxury_assets_value'
'bank_asset_value'
'loan_status'
```

```
In [32]: df.columns = df.columns.str.replace(" ", "")
```

```
In [ ]: # Mapping of categorical values of education, self_employed and loan_status to Numerical(0/1)
```

```
In [34]: df["education"] = df["education"].str.strip().map({"Graduate":1,"Not Graduate":0})
df["self_employed"] = df["self_employed"].str.strip().map({"Yes":1,"No":0})
df["loan_status"] = df["loan_status"].str.strip().map({"Approved":1,"Rejected":0})
df
```


Out[34]:

	loan_id	no_of_dependents	education	self_employed	income_annum	loan_amount	loan_term	cibil_score	residential_assets_value	commercial_assets_value	luxury_assets_value	bank_asset_value	loan_status
0	1	2	1	0	9600000	29900000	12	778	2400000	17600000	22700000	8000000	1
1	2	0	0	1	4100000	12200000	8	417	2700000	2200000	8800000	3300000	0
2	3	3	1	0	9100000	29700000	20	506	7100000	4500000	33300000	12800000	0
3	4	3	1	0	8200000	30700000	8	467	18200000	3300000	23300000	7900000	0
4	5	5	0	1	9800000	24200000	20	382	12400000	8200000	29400000	5000000	0
...
4264	4265	5	1	1	1000000	2300000	12	317	2800000	500000	3300000	800000	0
4265	4266	0	0	1	3300000	11300000	20	559	4200000	2900000	11000000	1900000	1
4266	4267	2	0	0	6500000	23900000	18	457	1200000	12400000	18100000	7300000	0
4267	4268	1	0	0	4100000	12800000	8	780	8200000	700000	14100000	5800000	1
4268	4269	1	1	0	9200000	29700000	10	607	17800000	11800000	35700000	12000000	1

4269 rows × 13 columns


In [35]:

```
gb = df.groupby("loan_status").agg({"loan_id":"count"})
plt.pie(gb["loan_id"],labels = gb.index,autopct = "%1.2f%%")
plt.show()
```




In [38]:


```
sns.countplot(x = df["self_employed"], data = df, hue = df["loan_status"])
plt.show()
```




```
In [40]: sns.boxplot(data=df, x= "loan_status",y= "cibil_score")
plt.show()
```



```
In [42]: sns.scatterplot(data=df, x= "income_annum", y="loan_amount", hue = "loan_status", palette= "magma")
plt.show()
```



```
In [44]: sns.countplot(data = df, x= "loan_term", hue = "loan_status", palette = "viridis")
plt.show()
```



```
In [46]: features = [
    "education", "no_of_dependents", "self_employed", "income_annum", "loan_amount", "loan_term", "cibil_score", "residential_assets_value", "commercial_assets_value", "bank_asset_value"
]
X = df[features].copy()
y = df["loan_status"]
```

```
In [47]: X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2, stratify=y, random_state=42)
```

```
In [48]: model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
```

Out[48]:

▼ LogisticRegression ⓘ ⓘ

► Parameters

```
In [49]: y_pred = model.predict(X_test)
y_proba = model.predict_proba(X_test)[:, 1]
```

```
In [50]: print("Accuracy : ", accuracy_score(y_test, y_pred)) #how many prediction were correct from the total
print("Precision:", precision_score(y_test, y_pred)) #all the positive predictions that models has made, from how many were actually positive
print("Recall : ", recall_score(y_test, y_pred)) #from the actual positives how many were model was able to identify
print("F1-score :", f1_score(y_test, y_pred)) #mean b/w recall and precision.
print("ROC-AUC : ", roc_auc_score(y_test, y_proba)) # 0 - 1, if towards 1, the model is good. if 0.5 (random guessing)
```

```
Accuracy : 0.8009367681498829
Precision: 0.7934959349593496
Recall    : 0.9190207156308852
F1-score  : 0.8516579406631762
ROC-AUC   : 0.8830351052106837
```

Model Performance Summary

- ◆ Accuracy: 0.80
 - The model predicts correctly 80% of the time.
- ◆ Precision: 0.79
 - When the model predicts loan approved, it is correct 79% of the time.
- ◆ Recall: 0.91
 - The model detects 91% of all actual approved loans.
- ◆ F1-Score: 0.85
 - Balanced measure of precision and recall.
- ◆ ROC-AUC: 0.88
 - Very good discrimination between approved (1) and rejected (0) customers.

```
In [51]: cm = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:\n", cm)
```

```
Confusion Matrix:
[[196 127]
 [ 43 488]]
```

```
In [56]: example = [
    1,          # education (0/1)
    2,          # no_of_dependents
    1,          # self_employed (0/1)
```

```

2700000,      # income_annum
15000000,     # loan_amount
15,           # loan_term
650,          # cibil_score
13500000,     # residential_assets_value
1200000,      # commercial_assets_value
1900000       # bank_asset_value
]]
pred_prob = model.predict_proba(example)[0, 1]
pred_class = model.predict(example)[0]

print("Predicted approval probability:", pred_prob)
print("Predicted class (1=Approved, 0=Rejected):", pred_class)

```

Predicted approval probability: 0.8647352492168883

Predicted class (1=Approved, 0=Rejected): 1

C:\Users\Girish Poojary\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\utils\validation.py:2749: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names

warnings.warn(

C:\Users\Girish Poojary\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\utils\validation.py:2749: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names

warnings.warn(

Executive Summary

- Overall system achieved strong predictive performance. With an accuracy of 80% and ROC-AUC of 0.88, the model demonstrates good ability to distinguish between high-risk and low-risk applicants.
- The analysis reveals that CIBIL score, income, asset values, and loan-to-income ratio are the strongest indicators of loan approval.
- Higher CIBIL scores particularly correlate with a significantly higher approval probability, confirming their importance in credit risk evaluation.
- The model is effective at identifying applicants who are genuinely eligible for loans, achieving a recall of 91%, meaning most good borrowers are accurately approved.

Overall, the model offers a reliable, data-driven decision-support tool that can improve approval consistency, reduce manual workload, and help minimize credit risk for financial institutions.

Key Insights

- ◆ 1. CIBIL Score is the Most Influential Factor Approved applicants show median CIBIL scores around 700+, while rejected applicants cluster near 400–450.
- ◆ 2. Higher Income and Higher Asset Values Increase Approval Odds Scatterplots indicate that applicants with higher annual income and higher residential or commercial asset values have a greater probability of approval.
- ◆ 3. Loan Term and Dependents Have Moderate Influence Shorter loan terms correlate with slightly higher approval rates. Applicants with fewer dependents show higher approval probability due to lower financial burden.
- ◆ 4. Model Performance is Strong but Can Improve Precision
- Confusion Matrix:

[[196 127]

[43 488]]

- True Positives (488): Approved applicants correctly identified.
- True Negatives (196): Rejected applicants correctly identified.
- False Positives (127): Risky applicants incorrectly approved.
- False Negatives (43): Eligible applicants incorrectly rejected.

Executive Recommendations

- ✓ 1. Use model for preliminary loan screening - High recall ensures that most legitimate borrowers are not rejected prematurely.
- ✓ 2. Strengthen rules for borderline approval cases - Use additional checks for applicants flagged near the decision threshold (0.5–0.7 probability).
- ✓ 3. Retrain model with more features - Adding debt-to-income ratio, employment history, and past default behavior can reduce false approvals.